

# Module7\_NeuralNets\_Tutorial

November 20, 2020

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
import pandas as pd

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

from sklearn.preprocessing import scale
from sklearn.metrics import mean_squared_error, roc_curve, auc
import statsmodels.formula.api as smf
from sklearn.model_selection import train_test_split
```

```
[2]: # Simulated Example I
```


```
[3]: def sigmoid(x):
    return(1 / (1 + np.exp(-x)))
```

```
[4]: np.random.seed(1)
x1 = np.random.normal(0, 1, 20000)
x2 = np.random.normal(0, 1, 20000)
y = (x1+0.5)*(x1+0.5)*(x1+0.5)*(x2-0.3)*(x2-0.3)+50*1/(1+np.exp(3*x1+2*x2))
```

```
[5]: fig = plt.figure(figsize=(16, 12))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, y, c='k', marker='o')

ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')


plt.show()
```



```
[6]: fig = plt.figure(figsize=(16, 12))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, y, c='k', marker='o')

ax.set_xlim3d(-2,2)
ax.set_ylim3d(-2,2)
ax.set_zlim3d(-20,80)
ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')

plt.show()
```



```
[7]: mydata1 = pd.DataFrame({'y':y, 'x1':x1, 'x2':x2})
```

```
[8]: mydata1_sc = scale(mydata1)
```

```
[9]: X_train = mydata1_sc[:,1:3]
     y_train = mydata1_sc[:,0]
```

```
[10]: inputs = keras.Input(shape=(2,))
      x = layers.Dense(2, activation="sigmoid", name="dense_1")(inputs)
      outputs = layers.Dense(1, activation="linear", name="predictions")(x)

      model = keras.Model(inputs=inputs, outputs=outputs)
```

```
[11]: model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.01),  # Optimizer
    # Loss function to minimize
    loss=keras.losses.MeanSquaredError(),
    # List of metrics to monitor
    metrics=[keras.metrics.MeanSquaredError()],
)
```

```
[12]: history = model.fit(
    X_train,
    y_train,
    # batch_size=64,
    epochs=5
)
```

Train on 20000 samples

Epoch 1/5

20000/20000 [=====] - 2s 95us/sample - loss: 0.3070 -  
mean\_squared\_error: 0.3070s - loss: 0.5575

Epoch 2/5

20000/20000 [=====] - 1s 61us/sample - loss: 0.1966 -  
mean\_squared\_error: 0.1966

Epoch 3/5

20000/20000 [=====] - 1s 50us/sample - loss: 0.1924 -  
mean\_squared\_error: 0.1924

Epoch 4/5

20000/20000 [=====] - 1s 53us/sample - loss: 0.1906 -  
mean\_squared\_error: 0.1906

Epoch 5/5

20000/20000 [=====] - 1s 51us/sample - loss: 0.1899 -  
mean\_squared\_error: 0.1899

```
[13]: pred_sc = model.predict(X_train)
preds_1 = pred_sc*np.std(mydata1).y + np.mean(mydata1).y
compare = pd.DataFrame({'y':y, 'preds_1':preds_1.T[0]})
compare.head()
```

```
[13]:      y    preds_1
0  61.091165  26.125568
1  44.811675  45.814064
2  48.131747  48.987610
3  48.661676  48.472984
4   1.816804   5.159172
```


```
[14]: mean_squared_error(mydata1.y, preds_1.T[0])
```

```
[14]: 73.26767286031229
```

```
[15]: fig = plt.figure(figsize=(16, 12))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, preds_1.T[0], c='k', marker='o')

ax.set_xlim3d(-2,2)
ax.set_ylim3d(-2,2)
ax.set_zlim3d(-20,80)
ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')

plt.show()
```




```
[16]: # Simulated Example II
```

```
[17]: np.random.seed(3)
x1 = np.random.normal(0, 1, 100)
x2 = np.random.normal(0, 1, 100)
x3 = np.random.normal(0, 1, 100)
y2_true = 3 + 2 * x1 + 4 * x2
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, y2_true, c='k', marker='o')

ax.set_xlim3d(-2,2)
ax.set_ylim3d(-2,2)
ax.set_zlim3d(-10,20)
ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')

plt.show()
```



```
[18]: y2 = y2_true + 5 * np.random.normal(0, 1, 100)
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, y2, c='k', marker='o')


ax.set_xlim3d(-2,2)
```

```

ax.set_ylim3d(-2,2)
ax.set_zlim3d(-20,30)
ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')

plt.show()

```



```

[19]: mydata2 = pd.DataFrame({'y2':y2, 'x1':x1, 'x2':x2, 'x3':x3})
mydata2_sc = scale(mydata2)

```

```

[ ]:

```

```

[20]: X_train = mydata2_sc[:,1:4]
y_train = mydata2_sc[:,0]

inputs = keras.Input(shape=(3,))
x = layers.Dense(5, activation="sigmoid", name="dense_1")(inputs)
x = layers.Dense(3, activation="sigmoid", name="dense_2")(x)
outputs = layers.Dense(1, activation="linear", name="predictions")(x)

model = keras.Model(inputs=inputs, outputs=outputs)
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.01), # Optimizer
    # Loss function to minimize
    loss=keras.losses.MeanSquaredError(),

```

```

    # List of metrics to monitor
    metrics=[keras.metrics.MeanSquaredError()],
)

history = model.fit(
    X_train,
    y_train,
    # batch_size=64,
    epochs=5
)

```

Train on 100 samples

```

Epoch 1/5
100/100 [=====] - 1s 10ms/sample - loss: 1.0204 -
mean_squared_error: 1.0204
Epoch 2/5
100/100 [=====] - 0s 332us/sample - loss: 0.9848 -
mean_squared_error: 0.9848
Epoch 3/5
100/100 [=====] - 0s 489us/sample - loss: 0.9667 -
mean_squared_error: 0.9667
Epoch 4/5
100/100 [=====] - 0s 340us/sample - loss: 0.9611 -
mean_squared_error: 0.9611
Epoch 5/5
100/100 [=====] - 0s 264us/sample - loss: 0.9708 -
mean_squared_error: 0.9708

```

```

[21]: pred2_sc = model.predict(X_train)
      preds_2 = pred2_sc*np.std(mydata2).y2 + np.mean(mydata2).y2

```

```


[22]: fig = plt.figure()
      ax = fig.add_subplot(111, projection='3d')
      ax.scatter(x1, x2, preds_2, c='k', marker='o')

      ax.set_xlim3d(-2,2)
      ax.set_ylim3d(-2,2)
      ax.set_zlim3d(-20,30)
      ax.set_xlabel('x1 Label')
      ax.set_ylabel('x2 Label')
      ax.set_zlabel('y Label')

      plt.show()

```






```
[23]: fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, y2_true-preds_2.T[0], c='k', marker='o')

ax.set_xlim3d(-2,2)
ax.set_ylim3d(-2,2)
ax.set_zlim3d(-20,30)
ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')

plt.show()
```



```
[24]: # single-layer networks with five neurons
```

```
[25]: inputs = keras.Input(shape=(3,))
x = layers.Dense(5, activation="sigmoid", name="dense_1")(inputs)
outputs = layers.Dense(1, activation="linear", name="predictions")(x)

model = keras.Model(inputs=inputs, outputs=outputs)
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.01), # Optimizer
    # Loss function to minimize
    loss=keras.losses.MeanSquaredError(),
    # List of metrics to monitor
    metrics=[keras.metrics.MeanSquaredError()],
)

history = model.fit(
    X_train,
    y_train,
    # batch_size=64,
    epochs=5
)
```

Train on 100 samples

Epoch 1/5

100/100 [=====] - 1s 7ms/sample - loss: 1.0854 -  
mean\_squared\_error: 1.0854

```

Epoch 2/5
100/100 [=====] - 0s 201us/sample - loss: 1.0101 -
mean_squared_error: 1.0101
Epoch 3/5
100/100 [=====] - 0s 267us/sample - loss: 0.9814 -
mean_squared_error: 0.9814
Epoch 4/5
100/100 [=====] - 0s 404us/sample - loss: 0.9651 -
mean_squared_error: 0.9651
Epoch 5/5
100/100 [=====] - 0s 316us/sample - loss: 0.9515 -
mean_squared_error: 0.9515

```

```

[26]: pred2_2_sc = model.predict(X_train)
      preds_2_2 = pred2_2_sc*np.std(mydata2).y2 + np.mean(mydata2).y2

```


```

[27]: fig = plt.figure()
      ax = fig.add_subplot(111, projection='3d')
      ax.scatter(x1, x2, y2_true-preds_2_2.T[0], c='k', marker='o')

      ax.set_xlim3d(-2,2)
      ax.set_ylim3d(-2,2)
      ax.set_zlim3d(-20,30)
      ax.set_xlabel('x1 Label')
      ax.set_ylabel('x2 Label')
      ax.set_zlabel('y Label')

      plt.show()

```



```
[28]: mean_squared_error(y2_true, preds_2_2.T[0])
```

```
[28]: 12.106177665098372
```

```
[29]: # With different hyperparameter: learning rate in this case

inputs = keras.Input(shape=(3,))
x = layers.Dense(5, activation="sigmoid", name="dense_1")(inputs)
outputs = layers.Dense(1, activation="linear", name="predictions")(x)

model = keras.Model(inputs=inputs, outputs=outputs)
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.05), # Optimizer
    # Loss function to minimize
    loss=keras.losses.MeanSquaredError(),
    # List of metrics to monitor
    metrics=[keras.metrics.MeanSquaredError()],
)

history = model.fit(
    X_train,
    y_train,
    # batch_size=64,
    epochs=5
)
```

Train on 100 samples

Epoch 1/5

100/100 [=====] - 1s 10ms/sample - loss: 1.0917 -  
mean\_squared\_error: 1.0917

Epoch 2/5

100/100 [=====] - 0s 161us/sample - loss: 0.8764 -  
mean\_squared\_error: 0.8764

Epoch 3/5

100/100 [=====] - 0s 148us/sample - loss: 0.8084 -  
mean\_squared\_error: 0.8084

Epoch 4/5

100/100 [=====] - 0s 171us/sample - loss: 0.7499 -  
mean\_squared\_error: 0.7499

Epoch 5/5

100/100 [=====] - 0s 125us/sample - loss: 0.7324 -  
mean\_squared\_error: 0.7324

```
[30]: pred2_3_sc = model.predict(X_train)
preds_2_3 = pred2_3_sc*np.std(mydata2).y2 + np.mean(mydata2).y2
```


```

fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')
ax.scatter(x1, x2, y2_true-preds_2_3.T[0], c='k', marker='o')

ax.set_xlim3d(-2,2)
ax.set_ylim3d(-2,2)
ax.set_zlim3d(-20,30)
ax.set_xlabel('x1 Label')
ax.set_ylabel('x2 Label')
ax.set_zlabel('y Label')

plt.show()

```



```
[31]: mean_squared_error(y2_true, preds_2_3.T[0])
```

```
[31]: 1.7573505065580433
```

```

[32]: lmfit = smf.ols(formula="y2 ~ x1 + x2 + x3", data=mydata2).fit()
yhat_OOS = lmfit.predict(mydata2.drop(columns = ["y2"]))
print(lmfit.summary())

```

#### OLS Regression Results

Dep. Variable:	y2	R-squared:	0.309
Model:	OLS	Adj. R-squared:	0.288
Method:	Least Squares	F-statistic:	14.33

Date: Fri, 20 Nov 2020 Prob (F-statistic): 8.68e-08  
Time: 14:24:45 Log-Likelihood: -308.23  
No. Observations: 100 AIC: 624.5  
Df Residuals: 96 BIC: 634.9  
Df Model: 3  
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.2953	0.561	4.089	0.000	1.181	3.410
x1	1.8233	0.524	3.481	0.001	0.784	2.863
x2	3.8645	0.634	6.096	0.000	2.606	5.123
x3	-0.7185	0.504	-1.426	0.157	-1.719	0.282
Omnibus:	0.466	Durbin-Watson:	2.037			
Prob(Omnibus):	0.792	Jarque-Bera (JB):	0.577			
Skew:	-0.150	Prob(JB):	0.749			
Kurtosis:	2.779	Cond. No.	1.52			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[33]: mean_squared_error(y2_true, yhat_OOS)
```

```
[33]: 1.4136322850403886
```

```
[34]: # Credit Card Defaults
```

```
[35]: mydata = pd.read_csv('UCI_Credit_Card.csv', index_col=0)
```

```
[36]: mydata.head()
```

```
[36]:
```

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	\
ID										
1	20000.0	2	2	1	24	2	2	-1	-1	
2	120000.0	2	2	2	26	-1	2	0	0	
3	90000.0	2	2	2	34	0	0	0	0	
4	50000.0	2	2	1	37	0	0	0	0	
5	50000.0	1	2	1	57	-1	0	-1	0	

	PAY_5	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	\
ID									
1	-2	...	0.0	0.0	0.0	0.0	689.0	0.0	
2	0	...	3272.0	3455.0	3261.0	0.0	1000.0	1000.0	
3	0	...	14331.0	14948.0	15549.0	1518.0	1500.0	1000.0	

```

4      0 ...    28314.0    28959.0    29547.0    2000.0    2019.0    1200.0
5      0 ...    20940.0    19146.0    19131.0    2000.0    36681.0    10000.0

```

```

      PAY_AMT4  PAY_AMT5  PAY_AMT6  default.payment.next.month
ID
1          0.0        0.0        0.0                        1
2       1000.0         0.0       2000.0                        1
3       1000.0      1000.0       5000.0                        0
4       1100.0      1069.0       1000.0                        0
5       9000.0        689.0        679.0                        0

```

[5 rows x 24 columns]

```
[37]: mydata.describe()
```

```

[37]:      LIMIT_BAL      SEX      EDUCATION      MARRIAGE      AGE \
count    30000.000000  30000.000000  30000.000000  30000.000000  30000.000000
mean     167484.322667      1.603733      1.853133      1.551867    35.485500
std      129747.661567      0.489129      0.790349      0.521970     9.217904
min       10000.000000      1.000000      0.000000      0.000000    21.000000
25%       50000.000000      1.000000      1.000000      1.000000    28.000000
50%      140000.000000      2.000000      2.000000      2.000000    34.000000
75%      240000.000000      2.000000      2.000000      2.000000    41.000000
max      1000000.000000      2.000000      6.000000      3.000000    79.000000

```

```

      PAY_0      PAY_2      PAY_3      PAY_4      PAY_5 \
count    30000.000000  30000.000000  30000.000000  30000.000000  30000.000000
mean       -0.016700     -0.133767     -0.166200     -0.220667     -0.266200
std         1.123802      1.197186      1.196868      1.169139      1.133187
min        -2.000000     -2.000000     -2.000000     -2.000000     -2.000000
25%        -1.000000     -1.000000     -1.000000     -1.000000     -1.000000
50%         0.000000      0.000000      0.000000      0.000000      0.000000
75%         0.000000      0.000000      0.000000      0.000000      0.000000
max         8.000000      8.000000      8.000000      8.000000      8.000000

```

```

      ...      BILL_AMT4      BILL_AMT5      BILL_AMT6      PAY_AMT1 \
count    ...    30000.000000  30000.000000  30000.000000  30000.000000
mean     ...    43262.948967  40311.400967  38871.760400    5663.580500
std      ...    64332.856134  60797.155770  59554.107537   16563.280354
min      ...   -170000.000000 -81334.000000 -339603.000000     0.000000
25%      ...     2326.750000    1763.000000    1256.000000    1000.000000
50%      ...    19052.000000   18104.500000   17071.000000    2100.000000
75%      ...    54506.000000   50190.500000   49198.250000    5006.000000
max      ...    891586.000000  927171.000000  961664.000000   873552.000000

```

```

      PAY_AMT2      PAY_AMT3      PAY_AMT4      PAY_AMT5 \
count    3.000000e+04  30000.000000  30000.000000  30000.000000

```

mean	5.921163e+03	5225.68150	4826.076867	4799.387633
std	2.304087e+04	17606.96147	15666.159744	15278.305679
min	0.000000e+00	0.000000	0.000000	0.000000
25%	8.330000e+02	390.000000	296.000000	252.500000
50%	2.009000e+03	1800.000000	1500.000000	1500.000000
75%	5.000000e+03	4505.000000	4013.250000	4031.500000
max	1.684259e+06	896040.000000	621000.000000	426529.000000

	PAY_AMT6	default.payment.next.month
count	30000.000000	30000.000000
mean	5215.502567	0.221200
std	17777.465775	0.415062
min	0.000000	0.000000
25%	117.750000	0.000000
50%	1500.000000	0.000000
75%	4000.000000	0.000000
max	528666.000000	1.000000

[8 rows x 24 columns]

```
[38]: factor = ['SEX', 'EDUCATION', 'MARRIAGE', 'PAY_0', 'PAY_2', 'PAY_3', 'PAY_4',
    ↪ 'PAY_5', 'PAY_6', 'default.payment.next.month']
mydata_numcols = mydata.drop(columns = factor)
mydata_faccols = mydata[factor].drop(columns = ['default.payment.next.month']).
    ↪ astype('category')
dummies = pd.get_dummies(mydata_faccols, drop_first=True)
mydata_numcols_sc_0 = scale(mydata_numcols)
mydata_numcols_sc = pd.DataFrame(data=mydata_numcols_sc_0, columns =
    ↪ mydata_numcols.columns, index = dummies.index)
mydata_sc = pd.concat([mydata_numcols_sc, dummies], axis = 1)
mydata_sc = pd.concat([mydata_sc, mydata['default.payment.next.month']], axis =
    ↪ 1)
```

```
[39]: train, test = train_test_split(mydata_sc, test_size=0.25)
X_train = train.drop(columns = ['default.payment.next.month']).values
y_train = train['default.payment.next.month'].values
X_test = test.drop(columns = ['default.payment.next.month']).values
y_test = test['default.payment.next.month'].values
```

```
[40]: inputs = keras.Input(shape=(82,))
x = layers.Dense(5, activation="relu", name="dense_1")(inputs)
x = layers.Dense(3, activation="sigmoid", name="dense_2")(x)
outputs = layers.Dense(1, activation="sigmoid", name="predictions")(x)

model = keras.Model(inputs=inputs, outputs=outputs)
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.01), # Optimizer
```



```

    # Loss function to minimize
    loss='binary_crossentropy',
    # List of metrics to monitor
    metrics=['accuracy'],
)

history = model.fit(
    X_train,
    y_train,
    # batch_size=64,
    epochs=15
)

```

Train on 22500 samples

Epoch 1/15

22500/22500 [=====] - 4s 174us/sample - loss: 0.4542 - accuracy: 0.8048

Epoch 2/15

22500/22500 [=====] - 2s 81us/sample - loss: 0.4399 - accuracy: 0.8172

Epoch 3/15

22500/22500 [=====] - 2s 80us/sample - loss: 0.4372 - accuracy: 0.8197

Epoch 4/15

22500/22500 [=====] - 2s 84us/sample - loss: 0.4351 - accuracy: 0.8194

Epoch 5/15

22500/22500 [=====] - 2s 85us/sample - loss: 0.4338 - accuracy: 0.8207

Epoch 6/15

22500/22500 [=====] - 2s 83us/sample - loss: 0.4330 - accuracy: 0.8192

Epoch 7/15

22500/22500 [=====] - 2s 80us/sample - loss: 0.4318 - accuracy: 0.8212

Epoch 8/15

22500/22500 [=====] - 2s 82us/sample - loss: 0.4315 - accuracy: 0.8206

Epoch 9/15

22500/22500 [=====] - 2s 80us/sample - loss: 0.4307 - accuracy: 0.8217

Epoch 10/15

22500/22500 [=====] - 2s 80us/sample - loss: 0.4301 - accuracy: 0.8204

Epoch 11/15

22500/22500 [=====] - 2s 83us/sample - loss: 0.4300 - accuracy: 0.8216


```
Epoch 12/15
22500/22500 [=====] - 2s 82us/sample - loss: 0.4293 -
accuracy: 0.8216
Epoch 13/15
22500/22500 [=====] - 2s 91us/sample - loss: 0.4285 -
accuracy: 0.8213
Epoch 14/15
22500/22500 [=====] - 2s 81us/sample - loss: 0.4285 -
accuracy: 0.8224
Epoch 15/15
22500/22500 [=====] - 3s 120us/sample - loss: 0.4283 -
accuracy: 0.8220
```

```
[41]: mypreds = model.predict(X_test)
```

```
[42]: mypreds_bin = mypreds > 0.5
table = pd.DataFrame({'True':y_test,'pred':mypreds_bin.T[0]})
table.groupby(['True','pred']).size().unstack('True')
```

```
[42]: True      0      1
pred
False  5624  1049
True   276   551
```

```
[43]: fpr, tpr, threshold = roc_curve(y_test, mypreds)
roc_auc = auc(fpr, tpr)
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



[ ]: